

A novel trading system for the stock market using Deep Q-Network action and instance selection

Myeongseok Park, Jaeyun Kim, David Enke

Soonchunhyang University, South Korea

<https://www.sciencedirect.com/science/article/abs/pii/S0957417424019109>

IIE4122 Midterm Project | 25.04.21 Mon

TEAM 13 | Kim Wonjun(2022147002), Kim Hyeonjin(2021125085), Jang Seohyun(2021190002)

Introduction

Consequently, trading systems designed to generate high stock market returns are developed via several supervised learning methods

- ✓ However, methods based on them make it difficult to adapt to the real-time nature of the stock market (can be noisy and fail to consider the nonlinear and complex nature of stock prices)

Contribution

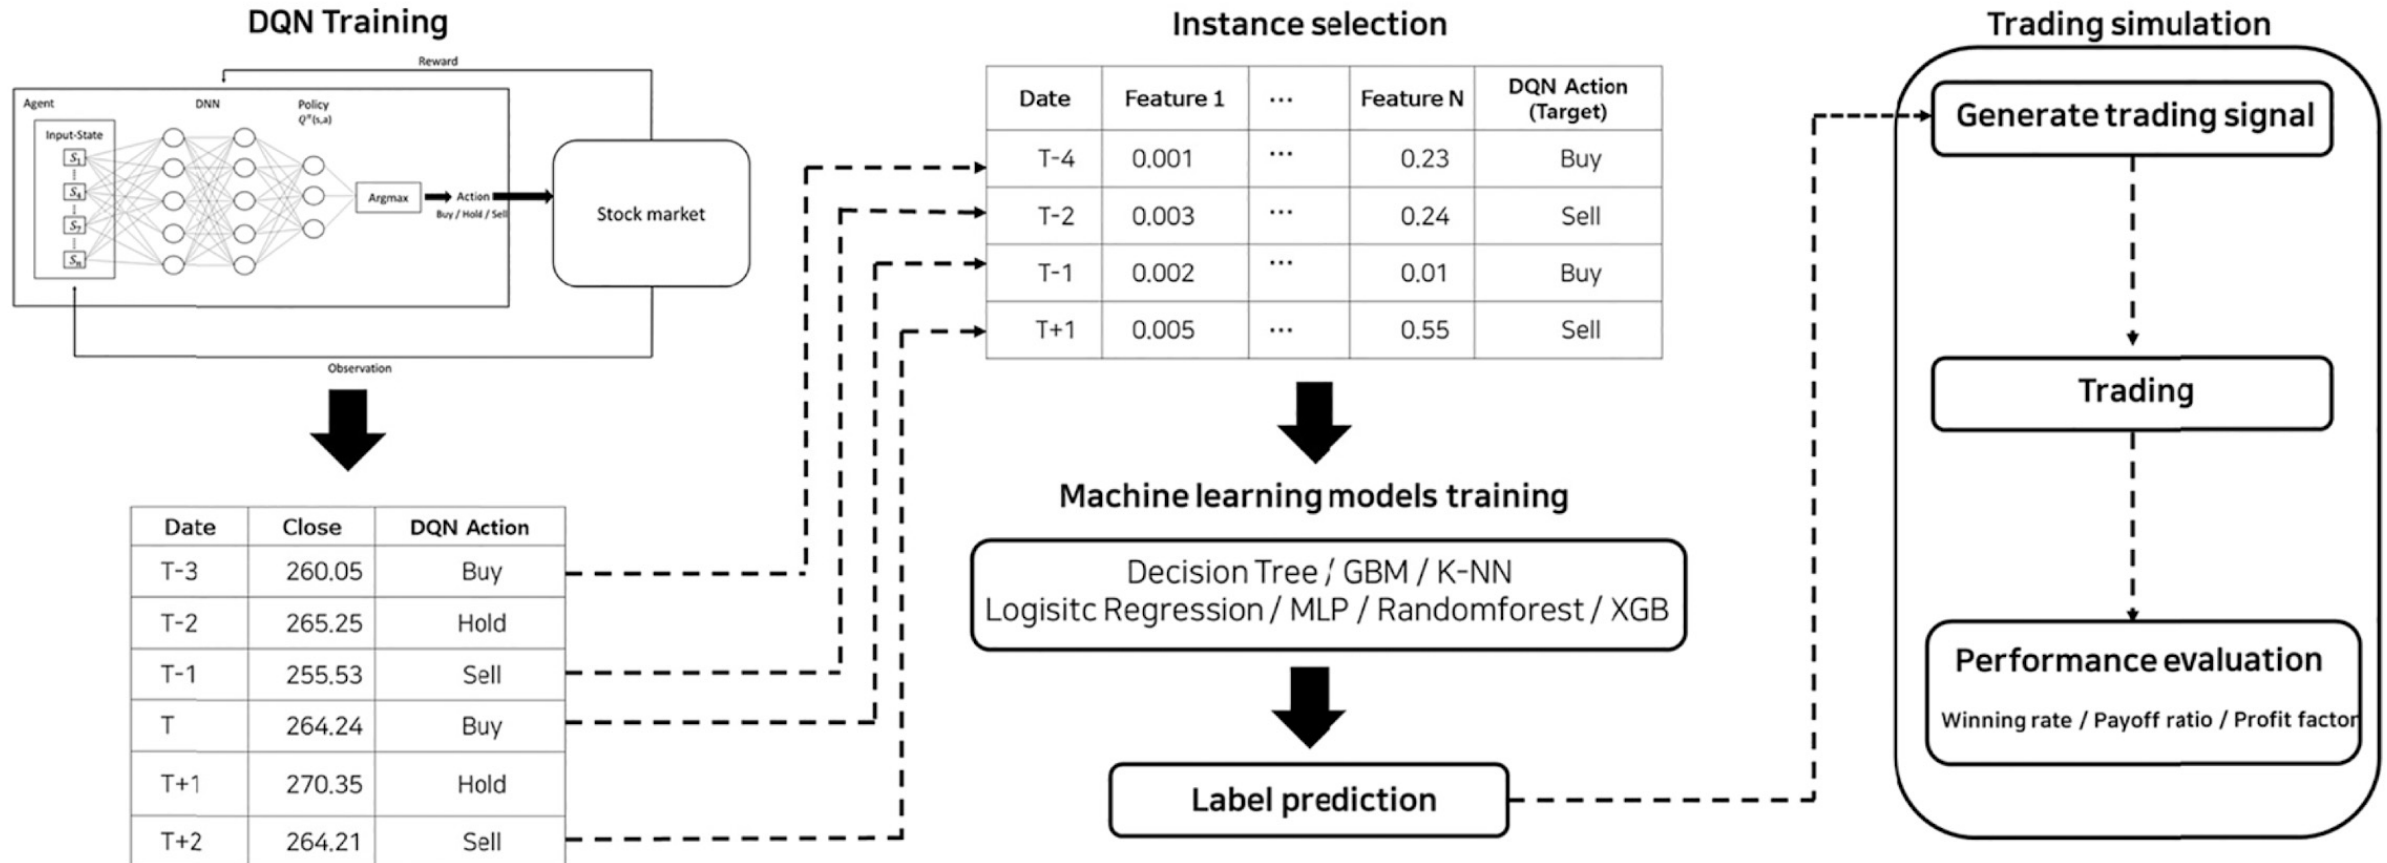
1. By applying reinforcement learning to stock trading, flexible strategies can be developed to adapt to changing market conditions
2. Existing literatures use price up/down labeling as labels, while this paper uses buy/sell signals trained by DQN



This study proposes a **Deep Q-Network (DQN) Action Instance Selection Trading System (DAIS)** to improve the limitations of both supervised learning and reinforcement learning trading systems

Methodology

2) The learning data (buy/sell signal + all technical indicators) and machine learning algorithm were constructed using instance selection




Contribution

✓ Limitation of Paper:

This study had some limitations. First, we arbitrarily selected the technical indicators to construct the reinforcement learning environment. The results may have been different if we had used more technical indicators or significant technical indicators. Second, we used the DQN

✓ Correlations among technical indicators



∴ Technical Indicator Selection is needed.

→ Feature selection is done using
1) Randomforest
2) mRMR

$N \rightarrow$ Select Top 20 features

Instance selection

Date	Feature 1	...	Feature N	DQN Action (Target)
T-4	0.001	...	0.23	Buy
T-2	0.003	...	0.24	Sell
T-1	0.002	...	0.01	Buy
T+1	0.005	...	0.55	Sell

Technical Indicator Selection

Machine learning models training

Decision Tree / GBM / K-NN
Logistic Regression / MLP / Randomforest / XGB

Label prediction

Improvement

Feature Selection code

```
import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from mrmr import mrmr_classif

feature_cols = [
    'ADX', 'ADXR', 'APO', 'aroondown', 'aroonup', 'AROONOSC', 'BOP', 'CCI', 'CMO', 'DX',
    'MACD', 'macdsignal', 'macdhist', 'ma_macd', 'ma_macdsignal', 'ma_macdhist',
    'fix_macd', 'fix_macdsignal', 'fix_macdhist', 'MFI', 'MINUS_DI', 'MINUS_DM',
    'MOM', 'PLUS_DM', 'PPO', 'ROC', 'ROCP', 'ROCR', 'ROCR100', 'RSI', 'slowk', 'slowd',
    'fastk', 'fastd', 'TRIX', 'ULTOSC', 'WILLR', 'two_crows', 'three_black_crows',
    'three_inside', 'three_line_strike', 'three_outside', 'three_star',
    'three_advance_white', 'abandoned_baby'
]

train_combined = []
for i in train_data:
    if 'label' in i.columns:
        tmp = i[feature_cols + ['label']].dropna()
        if not tmp.empty:
            train_combined.append(tmp)

data_all = pd.concat(train_combined, ignore_index=True)

X = data_all[feature_cols]
y = data_all['label']
y_bin = (y == 1).astype(int) # Buy = 1, Sell = 0

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_df = pd.DataFrame(X_scaled, columns=X.columns)
```

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_scaled, y_bin)
importances = rf.feature_importances_
selected_rf = X.columns[np.argsort(importances)[-20:]]

selected_mrmr = mrmr_classif(X=X_df, y=y_bin, K=20)

print("RandomForest Top 20 Features:", selected_rf.tolist())
print("MRMR Top 20 Features:", selected_mrmr)

import seaborn as sns
import matplotlib.pyplot as plt

corr_matrix = X_df.corr()
plt.figure(figsize=(15, 12))
sns.heatmap(corr_matrix, cmap='coolwarm', annot=False, fmt='.2f', linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.tight_layout()
plt.show()
```

Improvement

Selected Top 20 Features

RandomForest Top 20 Features: ['MOM', 'CCI', 'PPO', 'MINUS_DI', 'MACD', 'TRIX', 'fix_macd', 'ma_macdsignal', 'DX', 'MFI', 'ULTOSC', 'macdhist', 'ADXR', 'fix_macdhist', 'ma_macdhist', 'ADX', 'macdsignal', 'fix_macdsignal', 'MINUS_DM', 'PLUS_DM']

MRMR Top 20 Features: ['PLUS_DM', 'two_crows', 'ADX', 'MINUS_DI', 'MINUS_DM', 'ma_macdsignal', 'DX', 'ADXR', 'fix_macdsignal', 'ma_macdhist', 'macdsignal', 'BOP', 'fix_macd', 'MACD', 'APO', 'ma_macd', 'TRIX', 'macdhist', 'fix_macdhist', 'PPO']

→ We adopted top 20 features from **random forest**
according to the results of performance improvement

Results

Table 5
Trading performance by DQN and DAIS trading system.

Model	No. trades	Winning ratio	Payoff ratio	Profit factor	Sharpe ratio
Decision tree	90.30	0.48	1.10 (0.30)	1.11 (0.53)	1.03
GBM	88.27	0.47	1.10 (0.28)	1.10 (0.57)	0.66
knn	78.62	0.49	1.10 (0.40)	1.13 (0.67)	0.92
Logistic Regression	52.70	0.49	1.12 (0.59)	1.21 (1.93)	0.58
MLP	61.35	0.49	1.12 (0.53)	1.21 (0.89)	0.67
Random forest	76.98	0.49	1.15 (0.34)	1.20 (0.58)	1.09
XGB	89.04	0.49	1.04 (0.28)	1.09 (0.52)	0.93
DQN	84.01	0.46	0.97 (0.45)	0.94 (0.59)	0.38

Note: Values for trading performance are given as average (standard deviation).

Trading Performance of Original Paper

- ✓ Overall, the winning ratio has generally improved compared to the pre-improvement range (0.48~0.49).
- ✓ Several models achieved a winning ratio above 50%, suggesting a higher likelihood of making profitable trades.
- ✓ However, the number of trades has slightly decreased, averaging around 50–60 trades over two years (indicating relatively low trading frequency)
- ✓ Although the improved winning ratio implies better algorithmic performance, the small number of trades (n) limits the application of the law of large numbers, potentially introducing statistical noise.

methodology	year	stock_name	No.trades	Win%
pred_logistic	2019~2020	005930_삼성전자	33	0.666667
pred_decision	2019~2020	005930_삼성전자	98	0.581633
pred_naive (DQN)	2019~2020	005930_삼성전자	29	0.758621
pred_randomforest	2019~2020	005930_삼성전자	65	0.461538
pred_knn	2019~2020	005930_삼성전자	51	0.509804
pred_neural	2019~2020	005930_삼성전자	42	0.52381
pred_voting	2019~2020	005930_삼성전자	45	0.488889
pred_gbm	2019~2020	005930_삼성전자	107	0.336449

Trading Performance of Proposed Methodology (using RF)

Code Instruction Manual

- ✓ Development Environment: Python 3.10.1, Torch 2.5.1 => MUST DOWNGRADE
- ✓ TA-Lib must to be downloaded at <https://github.com/TA-Lib/ta-lib-python.git>
=> Then, run `!pip install ta-lib-0.6.4-src.tar.gz` within VScode (Linux version)

Windows

For 64-bit Windows, the easiest way is to get the *executable installer*:

1. Download [ta-lib-0.6.4-windows-x86_64.msi](#).
2. Run the Installer or run `msiexec` [from the command-line](#).

Alternatively, if you prefer to get the libraries without installing, or would like to use the 32-bit version:

- Intel/AMD 64-bit [ta-lib-0.6.4-windows-x86_64.zip](#)
- Intel/AMD 32-bit [ta-lib-0.6.4-windows-x86_32.zip](#)

Linux

Download [ta-lib-0.6.4-src.tar.gz](#) and:

```
$ tar -xzf ta-lib-0.6.4-src.tar.gz
$ cd ta-lib-0.6.4/
$ ./configure --prefix=/usr
$ make
$ sudo make install
```

If you build TA-Lib using `make -jX` it will fail but that's OK! Simply rerun `make -jX` followed by `[sudo] make install`.

- ✓ Other libraries(ex. ko_KR.UTF-8, mplfinance, tqdm) can be installed within VScode if needed (installation commands are inserted within submitted codes -> Google-colab basis)

Brief Instructions -> Detailed instructions are given

1. Run 'main.py'
 - 1a. Relocate prediction results to appropriate folders
2. Run 'main2.py'
3. Retrieve final results in .csv format

Code Instruction Manual

Run 'main.py' with the following commands sequentially

0. Modify path at line[7]

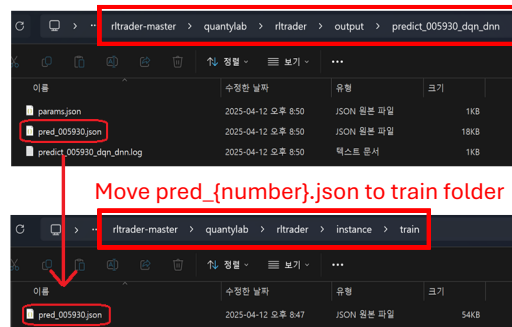
1. train data

```
python main.py --mode train --ver v5 --name 005930 --stock_code 005930 --rl_method dqn --net dnn --start_date 20130101 --end_date 20181231
```

2. predict train set

```
python main.py --mode predict --ver v5 --name 005930 --stock_code 005930 --rl_method dqn --net dnn --start_date 20130101 --end_date 20181231
```

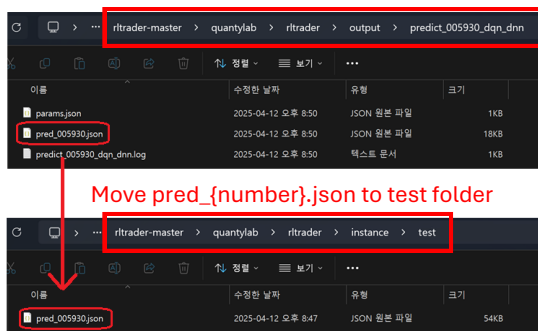
2a. Relocate train set prediction results



3. predict test set

```
python main.py --mode predict --ver v5 --name 005930 --stock_code 005930 --rl_method dqn --net dnn --start_date 20190101 --end_date 20201231
```

3a. Relocate test set prediction results



4. test data

```
python main.py --mode test --ver v5 --name 005930 --stock_code 005930 --rl_method dqn --net dnn --start_date 20190101 --end_date 20201231
```

Code Instruction Manual

Run 'main2.py' => Retrieve final results in .csv format

Example of final result:

year	stock_name	No.trades	Win%	Average gain	Average loss	Payoff ratio	Profit factor	Model
2019~2020	005930_삼	33	0.666667	3809.166	3489.998	1.091452	2.182905	pred_logistic
2019~2020	005930_삼	98	0.581633	2691.063	2805.99	0.959042	1.333303	pred_decision
2019~2020	005930_삼	29	0.758621	3203.988	3172.075	1.01006	3.174476	pred_naive
2019~2020	005930_삼	65	0.461538	1853.863	2035.293	0.910858	0.780736	pred_randomforest
2019~2020	005930_삼	51	0.509804	1050.141	1388.124	0.756518	0.786779	pred_knn
2019~2020	005930_삼	42	0.52381	1439.551	1464.789	0.98277	1.081048	pred_neural
2019~2020	005930_삼	45	0.488889	1347.668	1472.151	0.915441	0.87564	pred_voting
2019~2020	005930_삼	107	0.336449	1171.088	1673.365	0.69984	0.354849	pred_gbm

We have attached an inference code 'inference.ipynb' to help with code execution

Thank You

IIE4122 Midterm Project | 25.04.21 Mon

TEAM 13 | Kim Wonjun(2022147002), Kim Hyeonjin(2021125085), Jang Seohyun(2021190002)